


# Linux Observability and Tuning using bpftrace

## Module 4: Networking & Security Observability

Instructor: Chandrashekar Babu <training@chandrashekar.info>



# Agenda

## Network Tracing Fundamentals

Understanding socket operations, TCP/UDP tracing basics, and network stack instrumentation points

## Security Monitoring with bpftrace

Detecting suspicious behavior, unauthorized connection attempts, and monitoring process activities

## Advanced Network Observability

Deep packet inspection, NIC queue monitoring, and network performance analytics

## Hands-On Labs


Practical exercises implementing network and security monitoring solutions with bpftrace

# Network Tracing: Linux Networking Basics

Before diving into bpftrace, let's review where we can hook into the Linux networking stack:

- **Socket Layer:** High-level interface used by applications
- **Transport Layer:** TCP/UDP protocol implementations
- **Network Layer:** IP routing and forwarding
- **Link Layer:** Device drivers and physical transmission

bpftrace gives us visibility at all these layers, from socket operations down to device driver interactions.



# Key bpftrace Probe Points for Networking

## Socket Operations

Tracepoints:

- `sock:*`
- `syscall:socket*`
- `syscall:connect`
- `syscall:accept*`

## TCP Operations

Tracepoints:

- `tcp:*`
- `net:tcp_*`
- `kprobe:tcp_*`

## Packet Processing

Tracepoints:

- `net:netif_receive_skb`
- `net:net_dev_xmit`
- `kprobe:dev_queue_xmit`

These provide the foundation for our networking observability scripts.

# Basic Network Connection Monitoring

Let's start with a simple bpftrace one-liner to monitor new TCP connections:

```
#!/usr/bin/env bpftrace

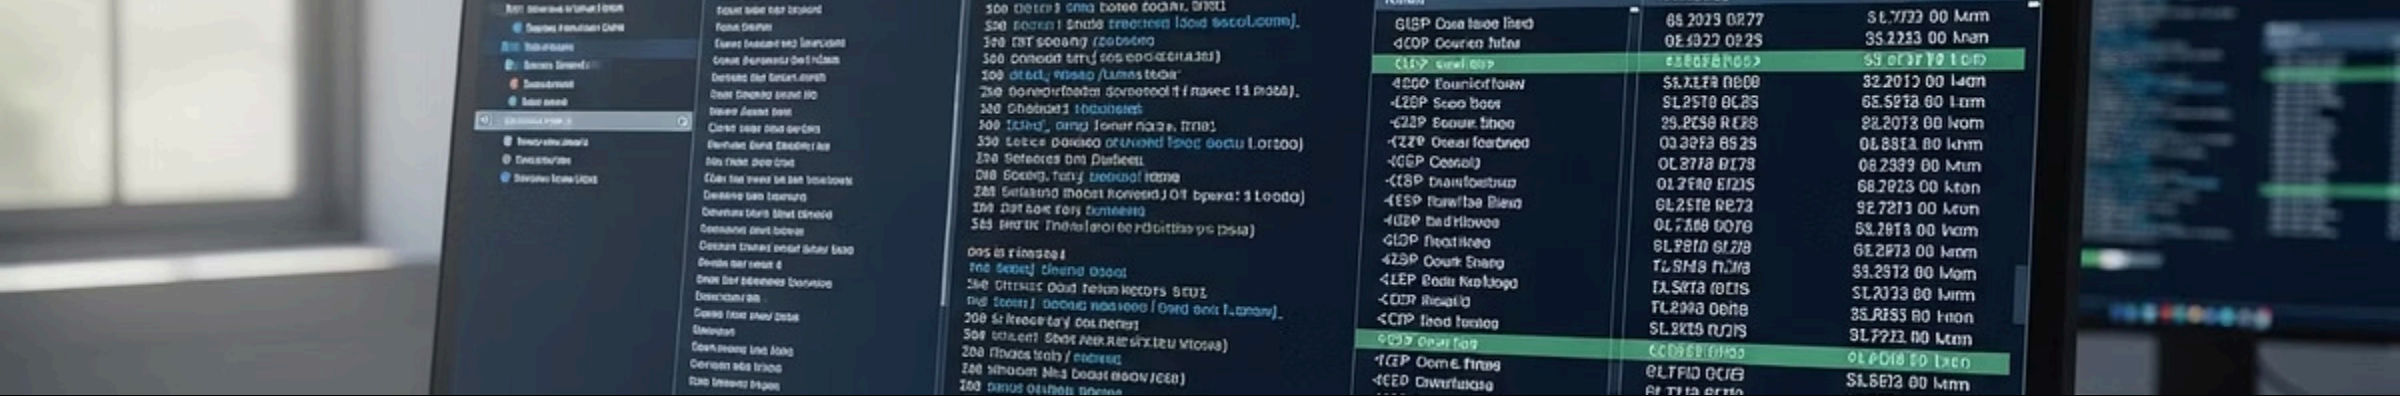
tracepoint:syscalls:sys_enter_connect
{
    $user_addr = arg2; // struct sockaddr * user_vaddr
    $family = *(uint16 *)($user_addr);

    if ($family == 2) { // AF_INET
        $port_be = *(uint16 *)($user_addr + 2);
        $ip = *(uint32 *)($user_addr + 4);

        // Convert network byte order to host byte order for port
        $port = (($port_be >> 8) & 0xFF) | (($port_be & 0xFF) << 8);

        printf("CONNECT: PID %d (%s) -> %s:%d\n",
            pid, comm,
            ntop($ip),
            $port);
    }
}
```

This one-liner tracks all connection attempts, showing the process name, PID, and destination IP/port.



# Tracking TCP Connection Lifetimes

The `tcplife.bt` script tracks the entire lifecycle of TCP connections:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("%-5s %-10s %-15s %-5s %-15s %-5s %s\n",
        "PID", "COMM", "LADDR", "LPORT", "RADDR", "RPORT", "MS");
}

kprobe:tcp_set_state
{
    $sk = (struct sock *)arg0;
    $newstate = arg1;

    if ($newstate == 1) {
        @start[$sk] = nsecs;
    }
}

kprobe:tcp_close
{
    $sk = (struct sock *)arg0;
    $delta = nsecs - @start[$sk];
    $durationms = $delta / 1000000;

    $family = $sk->__sk_common.skc_family;

    if ($family == 2) {
        $daddr = ntop($sk->__sk_common.skc_daddr);
        $saddr = ntop($sk->__sk_common.skc_rcv_saddr);
        $lport = $sk->__sk_common.skc_num;
        $dport = $sk->__sk_common.skc_dport;

        printf("%-5d %-10s %-15s %-5d %-15s %-5d %d\n",
            pid, comm, $saddr, $lport, $daddr, $dport, $durationms);
    }

    delete(@start[$sk]);
}
```

# Socket Connection Statistics

This one-liner counts active socket connections by process name:

```
bpfftrace -e 'tracepoint:syscalls:sys_enter_connect
{ @connects[comm] = count(); }
interval:s:5
{ print(@connects); clear(@connects); }'
```

For more detailed statistics including duration, use:

```
bpfftrace -e 'kprobe:tcp_connect { @start[arg0] = nsecs; }
kretprobe:tcp_connect /@start[arg0]/ {
  @duration_ns = hist(nsecs - @start[arg0]);
  delete(@start[arg0]);
}'
```

This shows a histogram of TCP connection durations, helping identify performance issues.

# Monitoring TCP Retransmits

TCP retransmissions can indicate network problems. Track them with:

```
bpfttrace -e 'kprobe:tcp_retransmit_skb {  
    @[pid, comm] = count();  
}'
```

For more context about what's being retransmitted:

```
bpfttrace -e 'kprobe:tcp_retransmit_skb {  
    $sk = (struct sock *)arg0;  
    $inet_family = $sk->__sk_common.skc_family;  
    if ($inet_family == AF_INET || $inet_family == AF_INET6) {  
        $daddr = ntop($sk->__sk_common.skc_daddr);  
        $saddr = ntop($sk->__sk_common.skc_rcv_saddr);  
        printf("Retransmit: %s -> %s (PID: %d, %s)\n",  
            $saddr, $daddr, pid, comm);  
    }  
}'
```



# Detailed TCP Retransmit Analysis

A more comprehensive script for detailed retransmit analysis:

```
#!/usr/bin/env bpftrace
```

```
BEGIN {
```

```
    printf("Tracing TCP retransmits... Hit Ctrl-C to end.\n");
```

```
    printf("%-8s %-8s %-16s %-16s %-5s %-5s %s\n",
```

```
    "TIME", "PID", "SOURCE", "DESTINATION", "SPORT", "DPORT", "STATE");
```

```
}
```

```
kprobe:tcp_retransmit_skb {
```

```
    $sk = (struct sock *)arg0;
```

```
    $inet_family = $sk->__sk_common.skc_family;
```

```
    if ($inet_family == AF_INET || $inet_family == AF_INET6) {
```

```
        $daddr = ntop($sk->__sk_common.skc_daddr);
```

```
        $saddr = ntop($sk->__sk_common.skc_rcv_saddr);
```

```
        $lport = $sk->__sk_common.skc_num;
```

```
        $dport = ntohs($sk->__sk_common.skc_dport);
```

```
        $state = $sk->__sk_common.skc_state;
```

```
        time("%H:%M:%S ");
```

```
        printf("%-8d %-16s %-16s %-5d %-5d %d\n",
```

```
        pid, $saddr, $daddr, $lport, $dport, $state);
```

```
        @retransmits[$saddr, $daddr] = count();
```

```
    }
```

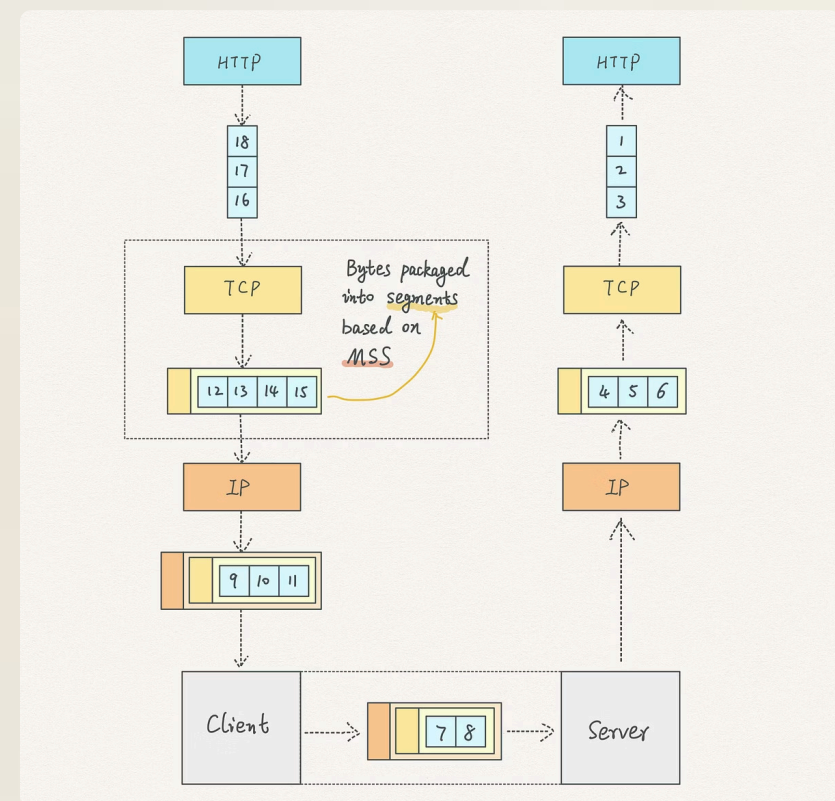
```
}
```

```
END {
```

```
    printf("\nRetransmit counts by connection:\n");
```

```
    print(@retransmits);
```

```
}
```



# Measuring TCP RTT (Round Trip Time)

This script measures TCP round-trip time by tracking ACK packets:

```
#!/usr/bin/env bpftrace

#include <linux/socket.h>
#include <net/inet_sock.h>

BEGIN {
    printf("Measuring TCP RTT...\n");
}

kprobe:tcp_transmit_skb {
    $sk = (struct sock *)arg0;
    $seq = $sk->tcp_sk.snd_nxt;
    @start[$sk, $seq] = nsecs;
}

kprobe:tcp_ack {
    $sk = (struct sock *)arg0;
    $seq = arg1;
    $start = @start[$sk, $seq];

    if ($start) {
        $rtt = (nsecs - $start) / 1000000; // ms

        // Only record if valid RTT
        if ($rtt > 0 && $rtt < 10000) {
            // Get connection details
            $daddr = ntop($sk->__sk_common.skc_daddr);
            $saddr = ntop($sk->__sk_common.skc_rcv_saddr);
            $lport = $sk->__sk_common.skc_num;
            $dport = ntohs($sk->__sk_common.skc_dport);

            printf("RTT: %6d ms (%s:%d -> %s:%d)\n",
                $rtt, $saddr, $lport, $daddr, $dport);

            // Record histogram by destination
            @rtt_hist[$daddr] = hist($rtt);
        }
        delete(@start[$sk, $seq]);
    }
}

interval:s:10 {
    print(@rtt_hist);
    clear(@rtt_hist);
}
```

# TCP SYN Flood Detection

Detect potential SYN flood attacks with this script:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("Monitoring for SYN flood attacks...\n");
    printf("Threshold: >100 SYN packets per second from a single source\n\n");
}

tracepoint:net:netif_receive_skb {
    $skb = (struct sk_buff *)args->skb;
    $ip = (struct iphdr *)($skb->head + $skb->network_header);

    // Check if it's an IPv4 packet
    if ($ip->version == 4) {
        // Extract the TCP header
        $tcp = (struct tcphdr *)($skb->head + $skb->transport_header);

        // Check if it's a SYN packet (SYN flag set, ACK flag not set)
        if (($tcp->syn == 1) && ($tcp->ack == 0)) {
            $saddr = ntop($ip->saddr);
            @syn_count[$saddr] = count();
        }
    }
}

interval:s:1 {
    // Check for potential SYN flood (high number of SYN packets from same source)
    printf("--- %s ---\n", strftime("%H:%M:%S", nsecs));

    foreach ($saddr in @syn_count) {
        $count = @syn_count[$saddr];
        if ($count > 100) {
            printf("ALERT: Possible SYN flood from %s (%d SYN packets/sec)\n",
                $saddr, $count);
        }
    }

    print(@syn_count);
    clear(@syn_count);
}
```

# Monitoring Network Socket Buffers

Track socket buffer usage to identify potential bottlenecks:

```
bpftrace -e 'kprobe:sock_alloc_send_skb {  
    @bytes[comm] = sum(arg1);  
}  
interval:s:5 {  
    printf("Socket send buffer allocation by process:\n");  
    print(@bytes);  
    clear(@bytes);  
}'
```

To monitor socket receive buffers:

```
bpftrace -e 'kprobe:__netif_receive_skb_core {  
    $skb = (struct sk_buff *)arg0;  
    @recv_bytes[cpu] = sum($skb->len);  
}  
interval:s:5 {  
    printf("Bytes received per CPU:\n");  
    print(@recv_bytes);  
    clear(@recv_bytes);  
}'
```

# DNS Query Monitoring

Track DNS queries on your system with this script:

```
#!/usr/bin/env bpftrace

#include <linux/socket.h>
#include <net/sock.h>
#include <linux/types.h>

BEGIN {
    printf("Tracing DNS queries... Hit Ctrl-C to end.\n");
}

// Catch UDP packets on port 53 (DNS)
kprobe:udp_sendmsg {
    $sk = (struct sock *)arg0;
    $dport = (uint16)($sk->__sk_common.skc_dport);

    // Check if it's to port 53 (DNS)
    if (ntohs($dport) == 53) {
        $daddr = ntop($sk->__sk_common.skc_daddr);
        printf("%-6d %-16s DNS query to %s (resolver)\n",
            pid, comm, $daddr);
        @dns_queries[comm] = count();
    }
}

END {
    printf("\nDNS query counts by process:\n");
    print(@dns_queries);
}
```

# Inspecting HTTP Request Headers

This advanced script detects and parses HTTP headers in network traffic:

```
#!/usr/bin/env bpftrace

#include <linux/socket.h>
#include <net/sock.h>
#include <linux/types.h>

BEGIN {
    printf("Monitoring HTTP requests... Hit Ctrl-C to end.\n");
    printf("%-6s %-16s %-16s %-5s %-5s %s\n",
        "PID", "COMM", "DSTIP", "PORT", "BYTES", "PATH");
}

kprobe:tcp_sendmsg {
    $sk = (struct sock *)arg0;
    $size = arg2;

    // Only process packets with a reasonable HTTP request size
    if ($size > 10 && $size < 1000) {
        // Get source buffer
        $iovbase = (char *)((struct iovec *)arg1)->iov_base;

        // Look for HTTP request pattern (GET, POST, PUT, etc.)
        $method = str($iovbase, 4);

        if (($method == "GET " || $method == "POST" ||
            $method == "PUT " || $method == "HEAD")) {

            // Extract destination IP and port
            $daddr = ntop($sk->__sk_common.skc_daddr);
            $dport = ntohs($sk->__sk_common.skc_dport);

            // If it's to standard HTTP/HTTPS ports
            if ($dport == 80 || $dport == 443 || $dport == 8080) {
                // Try to extract the request path
                $path_start = 0;
                $space_count = 0;

                // Find the second space which comes after the path
                for ($i = 0; $i < 80 && $i < $size; $i++) {
                    if ($iovbase[$i] == ' ') {
                        $space_count++;
                        if ($space_count == 1) {
                            $path_start = $i + 1;
                        } else if ($space_count == 2) {
                            // Found the end of the path
                            $path = str($iovbase + $path_start, $i - $path_start);
                            printf("%-6d %-16s %-16s %-5d %-5d %s\n",
                                pid, comm, $daddr, $dport, $size, $path);
                            break;
                        }
                    }
                }
            }

            @http_reqs[comm, $daddr] = count();
        }
    }
}

END {
    printf("\nHTTP request counts by process and destination:\n");
    print(@http_reqs);
}
```

# NIC Queue Monitoring

This one-liner tracks NIC transmit queue lengths:

```
bpfttrace -e 'kprobe:__netdev_pick_tx {
    @qlen[arg0->name] = hist(arg0->tx_queue_len);
}'
```

A more advanced script for monitoring NIC queue health:

```
#!/usr/bin/env bpfttrace

BEGIN {
    printf("Monitoring NIC queue metrics... Hit Ctrl-C to end.\n");
}

// Track queue fill level on transmit
kprobe:__dev_queue_xmit {
    $skb = (struct sk_buff *)arg0;
    $dev = $skb->dev;
    if ($dev) {
        @tx_qlen[$dev->name] = hist($dev->tx_queue_len);
        @tx_bytes[$dev->name] = sum($skb->len);
    }
}

// Track packet drops
kprobe:dev_hard_start_xmit {
    $ret = retval;
    if ($ret == -1) {
        $skb = (struct sk_buff *)arg0;
        $dev = $skb->dev;
        @drops[$dev->name] = count();
    }
}

interval:s:5 {
    time("%H:%M:%S\n");
    printf("NIC transmit queue lengths:\n");
    print(@tx_qlen);
    printf("\nBytes transmitted per interface:\n");
    print(@tx_bytes);
    printf("\nDropped packets per interface:\n");
    print(@drops);
    clear(@tx_qlen);
    clear(@tx_bytes);
}
```

# TCP Receive Window Analysis

This script monitors TCP receive window sizes to identify flow control issues:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("Analyzing TCP receive windows... Hit Ctrl-C to end.\n");
}

kprobe:tcp_rcv_established {
    $sk = (struct sock *)arg0;
    $tp = (struct tcp_sock *)$sk;

    // Calculate current receive window in bytes
    $rcv_wnd = $tp->rcv_wnd;
    $window_scaling = 1 << $tp->rx_opt.rcv_wscale;
    $effective_window = $rcv_wnd * $window_scaling;

    // Get connection details
    $daddr = ntop($sk->__sk_common.skc_daddr);
    $saddr = ntop($sk->__sk_common.skc_rcv_saddr);

    // Record histograms of receive window sizes
    @rcv_wnd_bytes[$daddr] = hist($effective_window);

    // Detect small windows (potential bottleneck)
    if ($effective_window < 4096) {
        printf("Small window alert: %s->%s window: %d bytes\n",
            $saddr, $daddr, $effective_window);
        @small_windows[$saddr, $daddr] = count();
    }
}

interval:s:10 {
    printf("\n=== TCP Receive Window Analysis ===\n");
    printf("Receive window size distributions by remote host:\n");
    print(@rcv_wnd_bytes);

    if (@small_windows) {
        printf("\nConnections with small receive windows:\n");
        print(@small_windows);
    }

    clear(@rcv_wnd_bytes);
    clear(@small_windows);
}
```



# Network Device Driver I/O

This script profiles network device driver activity:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("Tracing network device driver I/O... Hit Ctrl-C to end.\n");
}

// Trace network packet receipt
kprobe:netif_receive_skb {
    @receive_stack[kstack] = count();
    @dev_rx[arg0->dev->name] = count();
}

// Trace network packet transmission
kprobe:dev_hard_start_xmit {
    @transmit_stack[kstack] = count();
    @dev_tx[arg0->dev->name] = count();
}

interval:s:5 {
    printf("\n=== Network I/O by device ===\n");
    printf("Packets received:\n");
    print(@dev_rx);
    printf("\nPackets transmitted:\n");
    print(@dev_tx);

    clear(@dev_rx);
    clear(@dev_tx);
}

END {
    printf("\n=== Top packet receive stacks ===\n");
    print(@receive_stack, 5);
    printf("\n=== Top packet transmit stacks ===\n");
    print(@transmit_stack, 5);
}
```

# Socket Open & Close Tracking

Track socket lifecycle with this one-liner:

```
bpftool trace -e 'tracepoint:syscalls:sys_enter_socket { @opens[comm] = count(); }
tracepoint:syscalls:sys_enter_close /@socket_fds[tid]/ {
    @closes[comm] = count();
    delete(@socket_fds[tid]);
}
tracepoint:syscalls:sys_exit_socket /retval >= 0/ {
    @socket_fds[tid] = retval;
}'
```

For a more detailed view of socket operations by protocol:

```
bpftool trace -e 'tracepoint:syscalls:sys_enter_socket {
    @socket_by_type[arg0, arg1, arg2] = count();
}
interval:s:5 {
    printf("Socket creation by (domain, type, protocol):\n");
    print(@socket_by_type);
    clear(@socket_by_type);
}'
```

# Connection Tracking by Process

This comprehensive script tracks connections by process name:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("Tracing connection activity by process... Hit Ctrl-C to end.\n");
}

// Track socket creation
tracepoint:syscalls:sys_enter_socket {
    @socket_ops[pid, comm, "create"] = count();
}

// Track connection attempts
tracepoint:syscalls:sys_enter_connect {
    $sa = (struct sockaddr *)args->uservaddr;
    if ($sa->sa_family == AF_INET) {
        $in = (struct sockaddr_in *)$sa;
        $daddr = ntop($in->sin_addr.s_addr);
        $dport = ntohs($in->sin_port);

        printf("%-6d %-16s connect %s:%d\n", pid, comm, $daddr, $dport);
        @connect_to[comm, $daddr, $dport] = count();
        @socket_ops[pid, comm, "connect"] = count();
    }
}

// Track bind operations
tracepoint:syscalls:sys_enter_bind {
    $sa = (struct sockaddr *)args->umyaddr;
    if ($sa->sa_family == AF_INET) {
        $in = (struct sockaddr_in *)$sa;
        $addr = ntop($in->sin_addr.s_addr);
        $port = ntohs($in->sin_port);

        printf("%-6d %-16s bind %s:%d\n", pid, comm, $addr, $port);
        @bind_to[comm, $addr, $port] = count();
        @socket_ops[pid, comm, "bind"] = count();
    }
}

// Track listening sockets
tracepoint:syscalls:sys_enter_listen {
    @socket_ops[pid, comm, "listen"] = count();
}

// Track accept operations
tracepoint:syscalls:sys_enter_accept {
    @socket_ops[pid, comm, "accept"] = count();
}

// Track socket close
tracepoint:syscalls:sys_enter_close /@socket_fds[tid]/ {
    @socket_ops[pid, comm, "close"] = count();
    delete(@socket_fds[tid]);
}

tracepoint:syscalls:sys_exit_socket /args->ret >= 0/ {
    @socket_fds[tid] = args->ret;
}

interval:s:10 {
    printf("\n=== Socket operations by process ===\n");
    print(@socket_ops);

    printf("\n=== Connection attempts by destination ===\n");
    print(@connect_to);

    printf("\n=== Bind operations by address ===\n");
    print(@bind_to);

    clear(@socket_ops);
    clear(@connect_to);
    clear(@bind_to);
}
```

# Security Monitoring: Unauthorized Process Activity

This script detects suspicious network activity by unauthorized processes:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("Monitoring for unauthorized network activity...\n");

    // Define allowed network binaries - customize for your environment
    @allowed_net_bins["curl"] = 1;
    @allowed_net_bins["wget"] = 1;
    @allowed_net_bins["ssh"] = 1;
    @allowed_net_bins["scp"] = 1;
    @allowed_net_bins["rsync"] = 1;
    @allowed_net_bins["nc"] = 1;
    @allowed_net_bins["netcat"] = 1;
    @allowed_net_bins["firefox"] = 1;
    @allowed_net_bins["chrome"] = 1;
    @allowed_net_bins["chromium"] = 1;
}

tracepoint:syscalls:sys_enter_connect {
    $sa = (struct sockaddr *)args->uservaddr;

    // Only monitor IPv4 and IPv6 connections
    if ($sa->sa_family == AF_INET || $sa->sa_family == AF_INET6) {
        // Check if this is an allowed binary
        if (@allowed_net_bins[comm] != 1) {
            // This is a potential unauthorized connection
            if ($sa->sa_family == AF_INET) {
                $in = (struct sockaddr_in *)$sa;
                $daddr = ntop($in->sin_addr.s_addr);
                $dport = ntohs($in->sin_port);

                printf("ALERT: Unauthorized connection attempt by %s (PID %d): %s:%d\n",
                    comm, pid, $daddr, $dport);
                @unauth_conns[comm, $daddr, $dport] = count();

                // Record stack trace for investigation
                @stacks[comm, pid] = ustack;
            }
        }
    }
}

interval:s:10 {
    printf("\n=== Unauthorized connection summary ===\n");
    print(@unauth_conns);

    if (@stacks) {
        printf("\n=== Stack traces for investigation ===\n");
        print(@stacks);
        clear(@stacks);
    }
}
```

# Network Data Exfiltration Detection

Detect possible data exfiltration with large outbound transfers:

```
#!/usr/bin/env bpfttrace

BEGIN {
    printf("Monitoring for potential data exfiltration...\n");
    printf("Threshold: >10MB outbound in 1 minute\n");

    // Whitelist known data transfer processes - customize for your environment
    @allowed_data_xfer["rsync"] = 1;
    @allowed_data_xfer["scp"] = 1;
    @allowed_data_xfer["sftp"] = 1;
    @allowed_data_xfer["backup"] = 1;
}

kprobe:tcp_sendmsg {
    $sk = (struct sock *)arg0;
    $size = arg2;

    // Only track outbound connections (not localhost)
    if ($sk->__sk_common.skc_daddr != 0x0100007F && // 127.0.0.1
        $sk->__sk_common.skc_daddr != 0) {

        $daddr = ntop($sk->__sk_common.skc_daddr);
        $dport = ntohs($sk->__sk_common.skc_dport);

        // Track bytes by process
        @bytes[comm, $daddr, $dport] += $size;

        // Alert on large single writes
        if ($size > 1000000 && @allowed_data_xfer[comm] != 1) {
            printf("Large send: %s (PID %d) sending %d bytes to %s:%d\n",
                comm, pid, $size, $daddr, $dport);
        }
    }
}

interval:s:60 {
    printf("\n=== Checking for suspicious data transfers ===\n");

    // Detect potential exfiltration by checking total bytes transferred
    foreach ([ $comm, $daddr, $dport ] in @bytes) {
        $total = @bytes[ $comm, $daddr, $dport ];

        // Alert on large total transfers by non-whitelisted processes
        if ($total > 10000000 && @allowed_data_xfer[ $comm ] != 1) {
            printf("ALERT: Possible data exfiltration by %s to %s:%d (%d bytes)\n",
                $comm, $daddr, $dport, $total);
        }
    }

    printf("\n=== Data transfer summary (bytes) ===\n");
    print(@bytes);
    clear(@bytes);
}
```

# Port Scanning Detection

This script identifies potential port scanning activity:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("Monitoring for port scanning activity...\n");
    printf("Threshold: >15 different ports in 5 seconds from same source\n");
}

tracepoint:syscalls:sys_enter_connect {
    $sa = (struct sockaddr *)args->uervaddr;

    if ($sa->sa_family == AF_INET) {
        $in = (struct sockaddr_in *)$sa;
        $daddr = ntop($in->sin_addr.s_addr);
        $dport = ntohs($in->sin_port);

        // Track unique destinations by source process
        @dest_ports[pid, comm, $daddr, $dport] = count();

        // Count unique ports per destination
        @port_count[pid, comm, $daddr] = count();
    }
}

interval:s:5 {
    printf("\n=== Checking for port scan activity ===\n");

    // Check for processes connecting to many ports on the same host
    foreach ([ $pid, $comm, $daddr ] in @port_count) {
        $count = @port_count[$pid, $comm, $daddr];

        if ($count > 15) {
            printf("ALERT: Possible port scan by %s (PID %d) to %s (%d ports)\n",
                $comm, $pid, $daddr, $count);

            // Detailed breakdown of ports accessed
            printf("Ports accessed:\n");
            foreach ([ $pid2, $comm2, $daddr2, $dport ] in @dest_ports) {
                if ($pid2 == $pid && $daddr2 == $daddr) {
                    printf("  %d\n", $dport);
                }
            }
        }
    }

    // Only clear counters after alert generation
    clear(@dest_ports);
    clear(@port_count);
}
```

# Detecting Suspicious DNS Queries

This script monitors for DNS tunneling and other suspicious DNS activities:

```
#!/usr/bin/env bpftrace

#include
#include
#include
#include

BEGIN {
    printf("Monitoring for suspicious DNS activity...\n");
    printf("Checking for: long queries, high volume, and encoded data\n");
}

// Examine DNS packets (UDP port 53)
kprobe:udp_sendmsg {
    $sk = (struct sock *)arg0;
    $data = arg1;
    $size = arg2;
    $dport = ntohs($sk->__sk_common.skc_dport);

    // Only process DNS packets
    if ($dport == 53) {
        $daddr = ntop($sk->__sk_common.skc_daddr);

        // Count queries by process
        @dns_count[comm] = count();

        // Track total bytes (to detect data exfiltration)
        @dns_bytes[comm] += $size;

        // Basic checks for suspicious DNS activity
        if ($size > 100) {
            // Unusually large DNS query
            printf("Large DNS query (%d bytes) from %s (PID %d) to %s\n",
                $size, comm, pid, $daddr);
            @large_queries[comm, $daddr] = count();
        }
    }
}

interval:s:10 {
    printf("\n=== DNS Activity Analysis ===\n");

    // Check for high volume DNS queries (potential tunneling)
    foreach ($comm in @dns_count) {
        $count = @dns_count[$comm];
        $bytes = @dns_bytes[$comm];

        if ($count > 100) {
            printf("ALERT: High volume DNS queries from %s: %d queries, %d bytes\n",
                $comm, $count, $bytes);
        }
    }

    printf("\nDNS queries by process:\n");
    print(@dns_count);

    printf("\nLarge DNS queries (potential data encoding):\n");
    print(@large_queries);

    clear(@dns_count);
    clear(@dns_bytes);
}
```

# Process Network Isolation Monitoring

Track processes that should not be making network connections:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("Monitoring network activity by isolated processes...\n");

    // Define processes that should not make network connections
    // Add your specific processes here
    @isolated_procs["mysql"] = 1;
    @isolated_procs["postgres"] = 1;
    @isolated_procs["redis-server"] = 1;
}

tracepoint:syscalls:sys_enter_socket,
tracepoint:syscalls:sys_enter_connect,
tracepoint:syscalls:sys_enter_accept,
tracepoint:syscalls:sys_enter_sendto,
tracepoint:syscalls:sys_enter_recvfrom
{
    // Check if this is an isolated process
    if (@isolated_procs[comm] == 1) {
        printf("ALERT: Isolated process %s (PID %d) attempting network operation: %s\n",
            comm, pid, probe);
        @violations[comm, probe] = count();
        @stacks[comm, pid, probe] = ustack;
    }
}

interval:s:30 {
    if (@violations) {
        printf("\n=== Network isolation violations ===\n");
        print(@violations);

        printf("\n=== Stack traces for investigation ===\n");
        print(@stacks);
        clear(@stacks);
    }
}
```



# Socket Permission Violation Detection

Detect socket operations with insufficient permissions:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("Monitoring for socket permission violations...\n");
}

// Define a helper function to check uid/capability
// Note: this is simplified and should be enhanced in real scenarios
function check_socket_perm(prot, op) {
    // For demonstration - this should be customized for your environment

    // Check for privileged ports (<1024) being used by non-root
    if (((uint16)arg1 < 1024) && (uid != 0)) {
        printf("ALERT: Non-root user (UID %d) attempting to %s on privileged port %d\n",
            uid, op, (uint16)arg1);
        @priv_port_violations[comm, uid, (uint16)arg1, op] = count();
        return 1;
    }

    // Check for raw sockets (requires CAP_NET_RAW)
    if (prot == 3 && uid != 0) { // IPPROTO_RAW = 3
        printf("ALERT: Non-root user (UID %d) attempting to create raw socket\n", uid);
        @raw_socket_violations[comm, uid] = count();
        return 1;
    }

    return 0;
}

tracepoint:syscalls:sys_enter_socket {
    $domain = args->family;
    $type = args->type;
    $protocol = args->protocol;

    // Check permissions for socket creation
    check_socket_perm($protocol, "create socket");
}

tracepoint:syscalls:sys_enter_bind {
    $sa = (struct sockaddr *)args->umyaddr;

    // For IPv4 sockets
    if ($sa->sa_family == AF_INET) {
        $in = (struct sockaddr_in *)$sa;
        $port = ntohs($in->sin_port);

        // Check permissions for binding
        if ($port < 1024 && uid != 0) {
            printf("ALERT: Non-root user (UID %d) attempting to bind to privileged port %d\n",
                uid, $port);
            @bind_violations[comm, uid, $port] = count();
        }
    }
}

END {
    printf("\n=== Socket permission violation summary ===\n");
    printf("Privileged port violations:\n");
    print(@priv_port_violations);

    printf("\nRaw socket violations:\n");
    print(@raw_socket_violations);

    printf("\nBind violations:\n");
    print(@bind_violations);
}
```

# Monitoring File Descriptor Sharing

This script tracks file descriptor sharing between processes, which can include socket sharing:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("Monitoring for socket descriptor sharing between processes...\n");
}

// Track process relationships
tracepoint:syscalls:sys_enter_fork,
tracepoint:syscalls:sys_enter_vfork,
tracepoint:syscalls:sys_enter_clone {
    @parent[tid] = pid;
}

// Track socket creations
tracepoint:syscalls:sys_exit_socket /args->ret >= 0/ {
    $fd = args->ret;
    @socket_owner[$fd, pid] = comm;
}

// Track socket sharing via sendmsg with SCM_RIGHTS
kprobe:__sys_sendmsg {
    $msghdr = (struct msghdr *)arg1;
    $control = $msghdr->msg_control;

    // If control data exists, it might contain SCM_RIGHTS
    if ($control != 0) {
        // This is a simplification - actual SCM_RIGHTS detection would require
        // more detailed parsing of the cmsghdr structure
        printf("Potential FD sharing: %s (PID %d) sending control message\n",
            comm, pid);
        @potential_fd_sharing[comm, pid] = count();
    }
}

// Track when a process uses a socket it didn't create
kprobe:sock_sendmsg,
kprobe:sock_recvmsg {
    $sock = (struct socket *)arg0;
    $owner_pid = 0;
    $owner_found = 0;

    // Search for the real owner
    // This is simplified - in reality, we'd need to track FDs throughout their lifecycle
    foreach ([ $fd, $opid ] in @socket_owner) {
        if ($opid != pid) {
            $owner_found = 1;
            $owner_pid = $opid;
            $owner_comm = @socket_owner[$fd, $opid];

            printf("Socket sharing detected: %s (PID %d) using socket owned by %s (PID %d)\n",
                comm, pid, $owner_comm, $owner_pid);
            @shared_sockets[comm, pid, $owner_comm, $owner_pid] = count();
            break;
        }
    }

    interval:s:30 {
        printf("\n=== Socket sharing summary ===\n");

        if (@shared_sockets) {
            printf("Processes using sockets they didn't create:\n");
            print(@shared_sockets);
        }

        if (@potential_fd_sharing) {
            printf("\nPotential file descriptor sharing via SCM_RIGHTS:\n");
            print(@potential_fd_sharing);
        }
    }
}
```

# Network Connection Profiling

This script builds network connection profiles to detect anomalies:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("Building network connection profiles...\n");
    printf("Will alert on deviations from established patterns\n");
}

// Track connection patterns by process
tracepoint:syscalls:sys_enter_connect {
    $sa = (struct sockaddr *)args->servaddr;

    if ($sa->sa_family == AF_INET) {
        $in = (struct sockaddr_in *)$sa;
        $daddr = ntop($in->sin_addr.s_addr);
        $dport = ntohs($in->sin_port);

        // Build profile by process
        @conn_count[comm] = count();
        @dest_by_proc[comm, $daddr] = count();
        @port_by_proc[comm, $dport] = count();

        // Record first seen timestamp if this is a new destination
        if (@first_seen[comm, $daddr] == 0) {
            @first_seen[comm, $daddr] = nsecs;
        }
    }

    // Analyze data transfer patterns
    kprobe:tcp_sendmsg {
        $sk = (struct sock *)arg0;
        $size = arg2;

        if ($sk->__sk_common.skc_family == AF_INET) {
            $daddr = ntop($sk->__sk_common.skc_daddr);
            $dport = ntohs($sk->__sk_common.skc_dport);

            // Track data volumes
            @bytes_sent[comm, $daddr, $dport] += $size;

            // Detect sudden large transfers
            if (!@max_size[comm, $daddr, $dport]) {
                @max_size[comm, $daddr, $dport] = $size;
            } else if ($size > @max_size[comm, $daddr, $dport] * 5) {
                printf("Unusual large transfer: %s sending %d bytes to %s:%d (5x previous max)\n",
                    comm, $size, $daddr, $dport);
                @anomalies[comm, "large_transfer", $daddr, $dport] = count();
            } else if ($size > @max_size[comm, $daddr, $dport]) {
                @max_size[comm, $daddr, $dport] = $size;
            }
        }
    }

    interval:s:60 {
        printf("\n=== Network Connection Profile Analysis ===\n");

        // Calculate hourly rate for new connections
        $now = nsecs;

        foreach ([ $comm, $daddr ] in @first_seen) {
            $elapsed = ($now - @first_seen[$comm, $daddr]) / 1000000000; // seconds

            if ($elapsed < 3600 && @dest_by_proc[$comm, $daddr] > 10) {
                printf("New destination with high activity: %s connecting to %s (%d times in %.1f min)\n",
                    $comm, $daddr, @dest_by_proc[$comm, $daddr], $elapsed/60);
                @anomalies[$comm, "new_high_activity", $daddr] = count();
            }
        }

        // Report anomalies
        if (@anomalies) {
            printf("\nDetected network anomalies:\n");
            print(@anomalies);
            clear(@anomalies);
        }

        // Clean up old records - in real use this would be more sophisticated
        if (@conn_count) {
            clear(@first_seen);
        }
    }
}
```

# Socket Bufferbloat Detection

This script identifies socket buffer issues that can lead to latency:

```
#!/usr/bin/env bpftrace

BEGIN {
    printf("Monitoring for socket bufferbloat issues...\n");
}

// Track queued data in socket write buffers
kprobe:sock_alloc_send_pskb {
    $sk = (struct sock *)arg0;
    $size = arg1;

    if ($sk->__sk_common.skc_family == AF_INET) {
        $daddr = ntop($sk->__sk_common.skc_daddr);
        $dport = ntohs($sk->__sk_common.skc_dport);

        // Get current queue size
        $wmem_queued = $sk->sk_wmem_queued;

        // Record metrics
        @wmem_queued_hist[$daddr, $dport] = hist($wmem_queued);

        // Alert on large queues
        if ($wmem_queued > 1000000) { // 1MB threshold
            printf("Large socket write queue: %s -> %s:%d (%d bytes queued)\n",
                comm, $daddr, $dport, $wmem_queued);
            @large_queues[comm, $daddr, $dport] = $wmem_queued;
        }
    }
}

// Track receive buffer pressure
kprobe:tcp_rcv_established {
    $sk = (struct sock *)arg0;

    if ($sk->__sk_common.skc_family == AF_INET) {
        $saddr = ntop($sk->__sk_common.skc_rcv_saddr);
        $daddr = ntop($sk->__sk_common.skc_daddr);

        // Get receive buffer occupancy
        $rmem_alloc = $sk->sk_rmem_alloc;
        $rcvbuf = $sk->sk_rcvbuf;

        // Calculate percentage full
        $pct_full = ($rmem_alloc * 100) / $rcvbuf;

        // Record metrics
        @rmem_pct_hist[$saddr, $daddr] = hist($pct_full);

        // Alert on nearly full buffers
        if ($pct_full > 90) {
            printf("Receive buffer pressure: %s <- %s (%d%% full)\n",
                $saddr, $daddr, $pct_full);
            @rcv_pressure[$saddr, $daddr] = $pct_full;
        }
    }
}

interval:s:10 {
    printf("\n=== Socket Buffer Analysis ===\n");

    printf("Socket write buffer histograms (bytes queued):\n");
    print(@wmem_queued_hist);

    printf("\nSocket receive buffer fill percentage:\n");
    print(@rmem_pct_hist);

    if (@large_queues) {
        printf("\nLarge socket write queues (potential bufferbloat):\n");
        print(@large_queues);
        clear(@large_queues);
    }

    if (@rcv_pressure) {
        printf("\nReceive buffer pressure points:\n");
        print(@rcv_pressure);
        clear(@rcv_pressure);
    }

    clear(@wmem_queued_hist);
    clear(@rmem_pct_hist);
}
```

# Detecting Network Covert Channels

This advanced script detects potential covert channels in network traffic:

```
#!/usr/bin/env bpftrace

#include
#include
#include
#include

BEGIN {
    printf("Monitoring for network covert channels...\n");
    printf("Looking for: timing-based channels, uncommon protocol usage, header manipulation\n");
}

// Track TCP header field patterns
kprobe:tcp_sendmsg {
    $sk = (struct sock *)arg0;
    $tcp = (struct tcp_sock *)$sk;

    // Check for unusual TCP header options or values
    if ($tcp->tcp_header_len > 20) {
        // TCP header larger than standard 20 bytes indicates options
        @tcp_options_usage[comm, $tcp->tcp_header_len] = count();
    }

    // Check for abnormal sequence numbers or other patterns
    // This is a simplification - real detection would be more sophisticated
    if ($tcp->rcv_nxt == $tcp->copied_seq && $tcp->rcv_wup == $tcp->copied_seq) {
        @suspicious_seq[comm] = count();
    }
}

// Track uncommon protocol usage
tracepoint:syscalls:sys_enter_socket {
    $domain = args->family;
    $type = args->type;
    $protocol = args->protocol;

    // Look for rare protocol combinations
    if ($protocol > 10 && $protocol != 17 && $protocol != 6) {
        // Uncommon protocol number (not TCP or UDP)
        printf("Uncommon protocol: %s using protocol %d\n",
            comm, $protocol);
        @rare_protocols[comm, $protocol] = count();
    }
}

// Track packet timing patterns (potential timing channel)
kprobe:dev_hard_start_xmit {
    // Record timestamp of packet transmission
    @last_xmit[pid] = nsecs;
}

kprobe:dev_hard_start_xmit /@last_xmit[pid]/ {
    $interval = nsecs - @last_xmit[pid];

    // Look for suspicious timing patterns (e.g., very regular intervals)
    if ($interval > 0) {
        @xmit_intervals[comm] = hist($interval / 1000000); // ms

        // Check if previous 5 intervals were nearly identical
        // This is simplified - real detection would use more robust statistics
        if (@last_intervals[pid, 0] > 0 &&
            @last_intervals[pid, 1] > 0 &&
            @last_intervals[pid, 2] > 0 &&
            @last_intervals[pid, 3] > 0 &&
            @last_intervals[pid, 4] > 0) {

            $avg = (@last_intervals[pid, 0] +
                @last_intervals[pid, 1] +
                @last_intervals[pid, 2] +
                @last_intervals[pid, 3] +
                @last_intervals[pid, 4]) / 5;

            $dev0 = ($avg > @last_intervals[pid, 0]) ?
                ($avg - @last_intervals[pid, 0]) :
                (@last_intervals[pid, 0] - $avg);
            $dev1 = ($avg > @last_intervals[pid, 1]) ?
                ($avg - @last_intervals[pid, 1]) :
                (@last_intervals[pid, 1] - $avg);
            // ... (similar for dev2-4)

            // If all deviations are very small (< 1% of average)
            if ($dev0 < ($avg / 100) &&
                $dev1 < ($avg / 100)) {

                printf("Suspicious timing pattern detected: %s (PID %d) - regular intervals of ~%d ms\n",
                    comm, pid, $avg / 1000000);
                @timing_channels[comm, pid] = count();
            }
        }

        // Shift and store last 5 intervals
        @last_intervals[pid, 4] = @last_intervals[pid, 3];
        @last_intervals[pid, 3] = @last_intervals[pid, 2];
        @last_intervals[pid, 2] = @last_intervals[pid, 1];
        @last_intervals[pid, 1] = @last_intervals[pid, 0];
        @last_intervals[pid, 0] = $interval;
    }

    @last_xmit[pid] = nsecs;
}

interval:s:30 {
    printf("\n=== Covert Channel Detection Results ===\n");

    if (@tcp_options_usage) {
        printf("Processes using unusual TCP header options:\n");
        print(@tcp_options_usage);
    }

    if (@suspicious_seq) {
        printf("\nProcesses with suspicious TCP sequence patterns:\n");
        print(@suspicious_seq);
    }

    if (@rare_protocols) {
        printf("\nUncommon protocol usage:\n");
        print(@rare_protocols);
    }

    if (@timing_channels) {
        printf("\nPotential timing-based covert channels:\n");
        print(@timing_channels);
        clear(@timing_channels);
    }

    printf("\nPacket timing interval distributions:\n");
    print(@xmit_intervals);

    clear(@tcp_options_usage);
    clear(@suspicious_seq);
    clear(@xmit_intervals);
}
```





# Hands-On Lab: Unauthorized Connection Monitoring

In this lab, we'll develop and deploy a comprehensive tool to monitor for unauthorized connection attempts.

## Lab Overview:

1. Create a bpftrace script that identifies unauthorized connections
2. Enhance it to create detailed logs for forensic analysis
3. Integrate with a notification system for real-time alerts

# Lab Setup: Define Allowed Connections

First, let's create a baseline of allowed connections:

```
#!/usr/bin/env bpftool

BEGIN {
    printf("Unauthorized connection monitoring started...\n");

    // Define allowed connections by process and destination
    // Format: process_name:destination:port
    @allowed["nginx:10.0.0.5:3306"] = 1; // nginx to MySQL
    @allowed["java:10.0.0.6:27017"] = 1; // java app to MongoDB
    @allowed["python:10.0.0.7:6379"] = 1; // python app to Redis

    // Add your environment-specific rules here
}
```

Customize the allowed connections list to match your environment's legitimate network paths. This is critical for reducing false positives.

# Lab Part 1: Basic Connection Monitoring

Now let's add the core monitoring functionality:

```
// Add this to the previous script

// Track all connection attempts
tracepoint:syscalls:sys_enter_connect {
    $sa = (struct sockaddr *)args->uaddr;

    if ($sa->sa_family == AF_INET) {
        $in = (struct sockaddr_in *)$sa;
        $daddr = ntop($in->sin_addr.s_addr);
        $dport = ntohs($in->sin_port);

        // Create connection identifier string
        $conn_id = strcat(comm, ":");
        $conn_id = strcat($conn_id, $daddr);
        $conn_id = strcat($conn_id, ":");
        $conn_id = strcat($conn_id, str($dport));

        // Check if this connection is allowed
        if (@allowed[$conn_id] != 1) {
            printf("UNAUTHORIZED CONNECTION: %s (PID %d) -> %s:%d\n",
                comm, pid, $daddr, $dport);
            @unauth_conn[comm, $daddr, $dport] = count();
        }
    }
}
```



# Lab Part 2: Enhanced Forensic Logging

Let's extend our script with detailed forensic capabilities:

```
// Add this to the previous script

// Add user and process path information
tracepoint:syscalls:sys_enter_connect /@unauth_conn[comm, ntop(((struct sockaddr_in *)args->servaddr)->sin_addr.s_addr), ntohs(((struct
sockaddr_in *)args->servaddr)->sin_port)])/ {
    // Get user information
    $uid = uid;
    $gid = gid;

    // Log detailed information about this connection attempt
    printf("FORENSIC DETAIL - Time: %s\n", strftime("%H:%M:%S", nsecs));
    printf(" Process: %s (PID: %d, PPID: %d)\n", comm, pid, ppid);
    printf(" User: UID %d, GID %d\n", $uid, $gid);

    // Record stack trace for investigation
    printf(" User-space stack trace:\n");
    print(ustack);

    // Record kernel-space stack trace
    printf(" Kernel-space stack trace:\n");
    print(kstack);

    @detailed_events[comm, pid, $uid] = count();
}
```

# Lab Part 3: Implementing Real-Time Alerts

Now let's add real-time alerting functionality:

```
// Add this to the previous script

// Function to generate an alert message
function generate_alert(process, pid, uid, dst_ip, dst_port) {
    // In a real system, this would send to an external alerting system
    // For this lab, we'll simulate with a printf

    printf("\n!!! SECURITY ALERT !!!\n");
    printf("Unauthorized connection attempt detected:\n");
    printf(" Process: %s (PID: %d)\n", process, pid);
    printf(" User ID: %d\n", uid);
    printf(" Destination: %s:%d\n", dst_ip, dst_port);
    printf(" Timestamp: %s\n", strftime("%Y-%m-%d %H:%M:%S", nsecs));
    printf("!!! SECURITY ALERT !!!\n\n");
}

// Trigger alerts for unauthorized connections
tracepoint:syscalls:sys_enter_connect /@unauth_conn[comm, ntop(((struct sockaddr_in *)args->servaddr)->sin_addr.s_addr), ntohs(((struct
sockaddr_in *)args->servaddr)->sin_port)]/ {
    $sa = (struct sockaddr *)args->servaddr;

    if ($sa->sa_family == AF_INET) {
        $in = (struct sockaddr_in *)$sa;
        $daddr = ntop($in->sin_addr.s_addr);
        $dport = ntohs($in->sin_port);

        // Generate real-time alert
        generate_alert(comm, pid, uid, $daddr, $dport);
    }
}

// Summary report
interval:s:60 {
    printf("\n=== Unauthorized Connection Summary ===\n");
    time("%Y-%m-%d %H:%M:%S\n");

    if (@unauth_conn) {
        printf("Unauthorized connection attempts:\n");
        print(@unauth_conn);
    } else {
        printf("No unauthorized connections detected in the last minute\n");
    }

    clear(@unauth_conn);
}
```

# Lab Part 4: Running and Testing the Script

To run your unauthorized connection monitor:

```
# Save the complete script as unauthorized_conn_monitor.bt
# Make it executable
chmod +x unauthorized_conn_monitor.bt

# Run it with root privileges
sudo ./unauthorized_conn_monitor.bt

# Or use bpftrace directly
sudo bpftrace unauthorized_conn_monitor.bt
```

To test the script, try making connections from unauthorized processes:

```
# In another terminal, try an unauthorized connection
nc 10.0.0.5 3306

# You should see alerts in your bpftrace output
```

# Lab Part 5: Advanced Connection Fingerprinting

Let's enhance our monitor with connection fingerprinting:

```
// Add this feature to our monitoring script

// Track connection patterns
tracepoint:syscalls:sys_enter_connect {
    $sa = (struct sockaddr *)args->uaddr;

    if ($sa->sa_family == AF_INET) {
        $in = (struct sockaddr_in *)$sa;
        $daddr = ntop($in->sin_addr.s_addr);
        $dport = ntohs($in->sin_port);

        // Record connection timestamp
        @conn_time[pid, $daddr, $dport] = nsecs;

        // Count connections by process
        @conn_by_proc[comm] = count();
        @dest_by_proc[comm, $daddr] = count();
    }
}

// Track connection frequency patterns
interval:s:10 {
    printf("\n=== Connection Pattern Analysis ===\n");

    // Detect processes with unusual connection patterns
    foreach ($comm in @conn_by_proc) {
        $count = @conn_by_proc[$comm];

        if ($count > 20) {
            printf("High connection rate: %s made %d connections in 10 seconds\n",
                $comm, $count);
        }
    }

    // Detect processes connecting to many unique destinations
    $dest_count = 0;
    $prev_comm = "";

    foreach ([$comm, $daddr] in @dest_by_proc) {
        if ($prev_comm == $comm) {
            $dest_count++;
        } else {
            if ($dest_count > 10) {
                printf("Process connecting to many destinations: %s -> %d unique destinations\n",
                    $prev_comm, $dest_count);
            }
            $dest_count = 1;
            $prev_comm = $comm;
        }
    }

    // Clear counters for next interval
    clear(@conn_by_proc);
    clear(@dest_by_proc);
}
```

# Extending Our Security Observability

Beyond our lab, here are additional areas to explore for comprehensive security observability:

## Behavioral Baselineing

Create process-specific network behavior profiles and alert on deviations

## Traffic Correlation

Correlate network activity with system calls, file access, and memory patterns

## Lateral Movement Detection

Identify suspicious internal network traffic patterns that might indicate compromise

## Integration

Connect bpftape monitors with SIEM systems, log analyzers, and alerting platforms



# Key Takeaways: Network & Security Observability

## Network Observability

- bpftrace provides unprecedented visibility into network stack operations
- Trace from socket layer down to device driver for complete understanding
- Performance analysis can identify bottlenecks in packet processing
- Socket buffer monitoring helps detect queuing and latency issues

## Security Monitoring

- Real-time detection of unauthorized connections
- Identify suspicious network behavior patterns
- Monitor for covert channels and data exfiltration
- Track process network isolation violations

bpftrace empowers Linux systems engineers and security practitioners with deep network and security observability capabilities without requiring kernel modifications.

## Next Steps and Resources

## Additional Learning

- Brendan Gregg's "BPF Performance Tools" book
- Linux Tracing Workshops:  
<https://github.com/iovisor/bcc/blob/master/docs/tutorial.md>
- Networking Stack Internals: [Linux Networking Stack Guide](#)

# Build Your Own

- Extend the lab scripts for your specific environment
- Create a comprehensive network observability dashboard
- Implement custom bpftrace collectors for your monitoring system

Coming in Module 5

- Filesystem & Storage Observability
- Block I/O tracing
- Page cache analysis
- Storage performance monitoring